

Deep Learning (IST, 2022-23)

Practical 2: Perceptron

André Martins, Andreas Wichert, Luis Sá-Couto, Margarida Campos

Pen-and-Paper Exercises

The following questions should be solved by hand. You can use, of course, tools for auxiliary numerical computations.

Question 1

Consider the following linearly separable training set:

$$\mathbf{x}^{(1)} = \begin{bmatrix} -1 \\ 0 \end{bmatrix}, \mathbf{x}^{(2)} = \begin{bmatrix} 0 \\ 0.25 \end{bmatrix}, \mathbf{x}^{(3)} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, \mathbf{x}^{(4)} = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$
$$y^{(1)} = -1, y^{(2)} = +1, y^{(3)} = +1, y^{(4)} = -1.$$

1. Initialize all weights to zero (including the bias). Assume $\text{sign}(z) = +1$ iff $z \geq 0$, and -1 if $z < 0$. Use a learning rate of one. Compute two epochs of the perceptron learning algorithm.

Solution: We start by adding a constant feature of 1 to the inputs, leading to:

$$\mathbf{x}^{(1)} = \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix}, \mathbf{x}^{(2)} = \begin{bmatrix} 1 \\ 0 \\ 0.25 \end{bmatrix}, \mathbf{x}^{(3)} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \mathbf{x}^{(4)} = \begin{bmatrix} 1 \\ 1 \\ -1 \end{bmatrix}$$

The weight vector is initialized as $\mathbf{w} = [0, 0, 0]$, where the first dimension corresponds to the bias (intercept) parameter.

First epoch (3 mistakes in total):


- $\mathbf{w} = [0, 0, 0]^\top$
- $\hat{y}^{(1)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(1)}) = \text{sign}(0) = +1 \neq y^{(1)} - \text{mistake!}$
- Perceptron update: $\mathbf{w} \leftarrow \mathbf{w} + y^{(1)} \mathbf{x}^{(1)} = [0, 0, 0]^\top - [1, -1, 0]^\top = [-1, 1, 0]^\top$
- $\hat{y}^{(2)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(2)}) = \text{sign}(-1) = -1 \neq y^{(2)} - \text{mistake!}$
- Perceptron update: $\mathbf{w} \leftarrow \mathbf{w} + y^{(2)} \mathbf{x}^{(2)} = [-1, 1, 0]^\top + [1, 0, 0.25]^\top = [0, 1, 0.25]^\top$
- $\hat{y}^{(3)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(3)}) = \text{sign}(1.25) = +1 = y^{(3)} - \text{correct, no update.}$
- $\hat{y}^{(4)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(4)}) = \text{sign}(0.75) = +1 \neq y^{(4)} - \text{mistake!}$
- Perceptron update: $\mathbf{w} \leftarrow \mathbf{w} + y^{(4)} \mathbf{x}^{(4)} = [0, 1, 0.25]^\top - [1, 1, -1]^\top = [-1, 0, 1.25]^\top$

Second epoch (1 mistake in total):

- $\mathbf{w} = [-1, 0, 1.25]^\top$
- $\hat{y}^{(1)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(1)}) = \text{sign}(-1) = -1 = y^{(1)}$ – **correct, no update.**
- $\hat{y}^{(2)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(2)}) = \text{sign}(-0.6875) = -1 \neq y^{(2)}$ – **mistake!**
- Perceptron update: $\mathbf{w} \leftarrow \mathbf{w} + y^{(2)} \mathbf{x}^{(2)} = [-1, 0, 1.25]^\top + [1, 0, 0.25]^\top = [0, 0, 1.5]^\top$
- $\hat{y}^{(3)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(3)}) = \text{sign}(1.5) = +1 = y^{(3)}$ – **correct, no update.**
- $\hat{y}^{(4)} = \text{sign}(\mathbf{w}^\top \mathbf{x}^{(4)}) = \text{sign}(-1.5) = -1 = y^{(4)}$ – **correct, no update.**

2. After applying the algorithm until convergence the final trained weights were: $\mathbf{w} = [0, 1, 1.75]^\top$. Draw the separation hyperplane.

Solution: The trained weight vector is $\mathbf{w} = [0, 1, 1.75]^\top$, resulting in the following separating boundary, which corresponds to the line with equation $w_0 + w_1 x_1 + w_2 x_2 = 0$:



3. What is the perceptron output for the query point $\begin{bmatrix} 0 & 1 \end{bmatrix}^\top$?

Solution: We have $\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x}_{\text{new}}) = \text{sign}([0, 1, 1.75]^\top [1, 0, 1]) = \text{sign}(1.75) = +1$.

Question 2

The perceptron can learn a relatively large number of functions. In this exercise, we focus on simple logical functions.

1. Show graphically that a perceptron can learn the logical NOT function. Give an example with specific weights.

Solution: The NOT function receives as input a logical value $x \in \{-1, +1\}$ and outputs its logical negation $y \in \{-1, +1\}$. We can enumerate all possible inputs and their outputs:


- For $x = -1$ the output is $y = +1$
- For $x = +1$ the output is $y = -1$

To show that a perceptron can learn a given function we just need to show that all points that require a positive output (+1) can be separated from all points that require a negative output by an hyperplane.

Since we are working with 1-dimensional inputs, an hyperplane is, in this case, a point.

So, is there a point that accurately separates the points? Yes, any point between -1 and $+1$ will achieve this.

An example hyperplane is given by the weight -1 and bias 0 , leading to $\hat{y} = -x$.



2. Show graphically that a perceptron can learn the logical AND function for two inputs. Give an example with specific weights.

Solution: The AND function receives as input a pair of logical values $\mathbf{x} \in \mathbb{R}^2$ and outputs its logical conjunction $y \in \{-1, +1\}$. We can enumerate all possible inputs and their outputs:


- For $\mathbf{x} = \begin{bmatrix} -1 & -1 \end{bmatrix}^\top$ the output is $y = -1$
- For $\mathbf{x} = \begin{bmatrix} -1 & +1 \end{bmatrix}^\top$ the output is $y = -1$
- For $\mathbf{x} = \begin{bmatrix} +1 & -1 \end{bmatrix}^\top$ the output is $y = -1$
- For $\mathbf{x} = \begin{bmatrix} +1 & +1 \end{bmatrix}^\top$ the output is $y = +1$

To show that a perceptron can learn a given function we just need to show that all points that require a positive output (+1) can be separated from all points that require a negative output by an hyperplane.

Since we are working with 2-dimensional inputs, an hyperplane is, in this case, a line.

So, is there a weight vector that defines a boundary line that accurately separates the points?

Yes, for instance $\mathbf{w} = \begin{bmatrix} -1 & 1 & 1 \end{bmatrix}^\top$ achieves:



3. Show graphically that a perceptron can learn the logical OR function for two inputs. Give an example with specific weights.

Solution: The OR function receives as input a pair of logical values $\mathbf{x} \in \mathbb{R}^2$ and outputs its logical conjunction $y \in \{-1, +1\}$. We can enumerate all possible inputs and their outputs:


- For $\mathbf{x} = \begin{bmatrix} -1 & -1 \end{bmatrix}^\top$ the output is $y = -1$
- For $\mathbf{x} = \begin{bmatrix} -1 & +1 \end{bmatrix}^\top$ the output is $y = +1$
- For $\mathbf{x} = \begin{bmatrix} +1 & -1 \end{bmatrix}^\top$ the output is $y = +1$
- For $\mathbf{x} = \begin{bmatrix} +1 & +1 \end{bmatrix}^\top$ the output is $y = +1$

To show that a perceptron can learn a given function we just need to show that all points that require a positive output (+1) can be separated from all points that require a negative output by an hyperplane.

Since we are working with 2-dimensional inputs, an hyperplane is, in this case, a line.

So, is there a weight vector that defines a boundary line that accurately separates the points?

Yes, for instance $\mathbf{w} = \begin{bmatrix} 1 & 1 & 1 \end{bmatrix}^\top$ achieves:



4. Show graphically that a perceptron can not learn the logical XOR function for two inputs.

Solution: The XOR function receives as input a pair of logical values $\mathbf{x} \in \mathbb{R}^2$ and outputs its logical conjunction $y \in \{-1, +1\}$. We can enumerate all possible inputs and their outputs:


- For $\mathbf{x} = \begin{bmatrix} -1 & -1 \end{bmatrix}^\top$ the output is $y = -1$
- For $\mathbf{x} = \begin{bmatrix} -1 & +1 \end{bmatrix}^\top$ the output is $y = +1$
- For $\mathbf{x} = \begin{bmatrix} +1 & -1 \end{bmatrix}^\top$ the output is $y = +1$
- For $\mathbf{x} = \begin{bmatrix} +1 & +1 \end{bmatrix}^\top$ the output is $y = -1$

To show that a perceptron can learn a given function we just need to show that all points that require a positive output (+1) can be separated from all points that require a negative output by an hyperplane.

Since we are working with 2-dimensional inputs, an hyperplane is, in this case, a line.

So, is there a weight vector that defines a boundary line that accurately separates the points?

No... If we plot the points we see right away that no line can separate the positive instances from the negative ones:



Programming Exercises

The following exercises should be solved using Python, you can use the corresponding practical's notebook for guidance.

1. Let us consider **Question 1**.
 - (a) Initialize all weights to zero (including the bias). Assume $\text{sign}(z) = +1$ iff $z \geq 0$, and -1 if $z < 0$. Use a learning rate of one. Apply the perceptron learning algorithm until convergence. How many epochs does it take to converge?
 - (b) Change the initialization of weights and biases to be random with a standard normal distribution $\mathcal{N}(0, 1)$. Try multiple times. Does it always converge?
2. Generate a balanced dataset with 30 examples in \mathbb{R}^2 and 3 classes. Assume each of the 10 inputs associated to class $k \in \{0, 1, 2\}$ is generated as $x \sim \mathcal{N}(\mu_k, \sigma_k^2 I)$, with $\sigma_0 = \sigma_1 = \sigma_2 = 1$, $\mu_0 = [0, 0]^\top$, $\mu_1 = [0, 3]^\top$, and $\mu_2 = [2, 2]^\top$. Plot the data.

- (a) Implement the multi-class perceptron algorithm and run 100 iterations. Initialize all the weights to zero and use a learning rate of one. What is the training accuracy (fraction of points that are correctly classified)?
- 3. Now it's time to try the perceptron on real data and see what happens. Load the UCI handwritten digits dataset using `scikit-learn`
 - (a) Randomly split this data into training (80%) and test (20%) partitions.
 - (b) Create and run your implementation of the multi-class perceptron algorithm on this dataset. Measure the training and test accuracy.
 - (c) Use `scikit-learn`'s implementation of the perceptron algorithm. Compare the resulting accuracies.